

# Sterling Suggs

swsuggs@student.byu.edu · +1 209 843 0298 · swsuggs.github.io

## Education

---

- **M.S. Computer Science, GPA: 3.93/4**  
Brigham Young University  
ML emphasis  
Thesis: *Reinforcement Learning with Memory Networks*  
2021
- **B.S. Applied Mathematics, GPA: 3.91/4**  
Brigham Young University  
CS minor  
2017

## Experience

---

- **BYU Computer Science Dept.**  
Research Assistant – Provo, UT  
- Researched applications of auxiliary memory in deep reinforcement learning, in areas of sample efficiency, representation learning, and meta-learning  
- Designed and implemented three memory models with read and write functions to efficiently query and update a huge network layer  
- Devised experiments to test models' ability to adapt to new data and learn faster than pure gradient propagation  
*September 2017 onward*
- **Veracity Forecasting and Analysis**  
Software Engineer – Brigham City, UT  
- Codeveloped a simulation of the Navy Enlistment Supply Chain which the Navy will use to identify and relieve bottlenecks  
- Programmed physics-based model of solid rocket motor fuel to improve prediction of failure  
- Created application to automate data collection for a human resources project  
*April 2016 to April 2017*
- **BYU Mathematics Dept.**  
Developer – Provo, UT  
- Drafted and refined new scientific computing and data science programming labs for university curricula  
- Wrote new course material to teach math majors Python programming  
*October 2015 to April 2016*

## Skills

---

- **Technologies**  
Python, C++, Java, Pytorch, Tensorflow, Pandas, Scipy/Scikit-learn, MongoDB, LaTeX, SQL
- **Platforms and Project Management**  
Linux, Docker, Slurm, Git, Mercurial, AWS, Bash

## Papers implemented

---

- **AlphaGo Zero** [deepmind.com]  
Reinforcement learning meets tree search to learn a two-player game through self-play
- **DQN** [deepmind.com]  
A deep neural network achieves superhuman performance on Atari video games
- **True Online TD( $\lambda$ )** [proceedings.mlr]  
Classic TD( $\lambda$ ) is improved with an algorithm that allows more frequent value updates
- **Linear High-D Interpolation** [ams.org]  
A simplex method sidesteps the curse of dimensionality to interpolate efficiently in high dimensions

## Extracurricular books

---

- *Reinforcement Learning: An Introduction*, Sutton and Barto, 2018
- *Machine Learning*, Kevin Murphy, 2012
- *Deep Learning*, Goodfellow, Bengio, Courville, 2015
- *The Elements of Statistical Learning*, Hastie, Tibshirani, Friedman, 2001

## Relevant coursework

---

- **Math**  
Algorithm Design and Optimization, Control Theory, Bayesian Statistics, PDEs, Mathematical Analysis
- **Computer Science**  
Multi-agent systems, Computer Vision, Deep Neural Networks, Computational Creativity